BiTempQA: A Diagnostic Benchmark for Bitemporal Reasoning in LLM Agent Memory Systems

Anonymous

Abstract

Memory-augmented LLM agents must track not only *what* happened but *when* it happened and *when they learned about it*—a dual-timestamp (bitemporal) distinction critical for real-world applications such as medical records, financial systems, and news analysis. Existing benchmarks either ignore timestamps entirely or treat time as a single dimension, leaving a gap in evaluating this capability.

We introduce **BiTempQA**, a diagnostic benchmark of 308 Chinese question-answer pairs across 10 scenario types and 9 question types at three difficulty levels, where each memory entry carries explicit `event_time` (when something occurred) and `record_time` (when the system stored it) annotations. Over half (56.5%) of questions require reasoning about both timestamps simultaneously. We evaluate five memory backends (FAISS, BM25, Naive RAG, Simple KG, ChromaDB) through a unified Retrieve→LLM Generate→LLM Judge pipeline. A mixed-effects logistic regression confirms that record-time reasoning is genuinely harder than event-time reasoning (p=0.023), surviving controls for question type, difficulty, and scenario variation. We further find that: (1) the top three systems are statistically tied overall (p>0.3) but diverge on specific question types (40.6pp accuracy spread); (2) an oracle ensemble of all systems achieves 85.4% (+9.8pp over any single system); (3) approximately 55% of failures stem from retrieval limitations rather than reasoning errors; and (4) system confidence scores are *anti-calibrated* (AUROC < 0.5), predicting failure better than success. Cross-validation with the LoCoMo benchmark reveals striking system asymmetries: Simple KG dominates adversarial questions (97.5%) but trails on temporal reasoning (37.5%). BiTempQA is released as an open diagnostic tool for developing temporally-aware memory systems.

1 Introduction

Consider a medical AI assistant that learns at 10:00 AM on Tuesday that a patient's symptoms began two days prior. To correctly answer "What was the patient's status on Sunday morning?", the system must distinguish *when the symptoms started* (the **event time**) from *when it recorded this fact* (the **record time**). This bitemporal distinction—well-established in temporal database theory (Snodgrass, 1999; Jensen and Snodgrass, 1999)—is pervasive: financial trades execute Monday but settle Wednesday, news events occur Thursday evening but are reported Friday morning, scientific findings are established in 2022 but published in 2024. A memory system that conflates these two timestamps will make incorrect temporal inferences.

Despite this practical importance, no existing benchmark systematically evaluates whether LLM agent memory systems can reason about both dimensions simultaneously. Temporal reasoning benchmarks (e.g., TIME (Zhou et al., 2025), ETRQA (Zhang et al., 2025), TEST OF TIME (Kashyap et al., 2025)) focus on LLM reasoning over static text, not over dynamically evolving memory stores. Memory system benchmarks (e.g., LoCoMo (Raji et al., 2024), LongMemEval (Wu et al., 2024), MemoryAgentBench (Wang et al., 2025)) include temporal questions but treat time as a single flat dimension, without separating event occurrence from information acquisition. This gap means we cannot diagnose whether current memory backends fail on temporal questions because of *retrieval limitations* or *reasoning limitations*, nor identify which types of temporal reasoning are hardest.

We present **BiTempQA**, a diagnostic benchmark that fills this gap. BiTempQA is designed around a core insight from bitemporal data modeling: every fact has two independent temporal dimensions. Our benchmark contains **308 Chinese QA pairs**

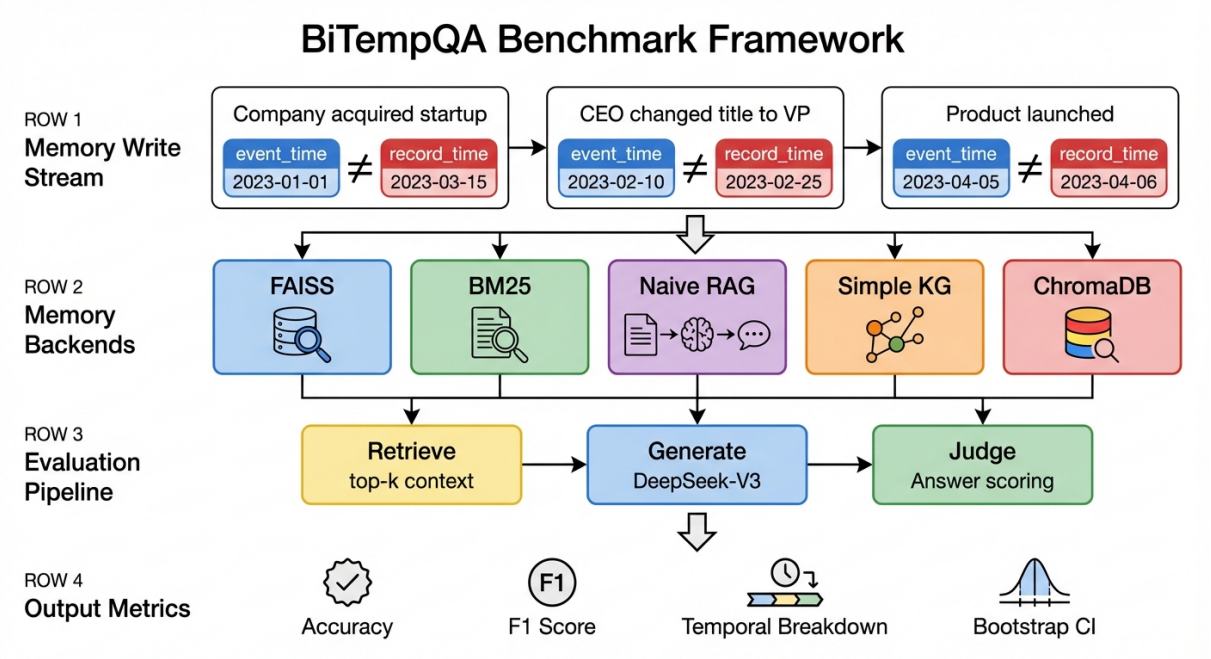

Figure 1: BiTempQA framework overview. Each memory write carries dual timestamps: `event_time` (when the event occurred) and `record_time` (when the system recorded it). Five memory backends retrieve context, which flows through a unified Retrieve→Generate→Judge evaluation pipeline to produce diagnostic metrics.

spanning **10 scenario types** across 6 domains (corporate, academic, social, fictional, historical, scientific), with **9 question types** at **3 difficulty levels**. Each memory write carries explicit `event_time` and `record_time` annotations, and 56.5% of questions require reasoning about *both* timestamps simultaneously. We evaluate five memory backends—FAISS, BM25, Naive RAG, Simple Knowledge Graph, and ChromaDB—through a unified Retrieve→Generate→Judge pipeline.

Our contributions are:

1. **First dual-timestamp diagnostic benchmark** for memory systems, with explicit event/record time annotations on every memory write and structured metadata on every question.

2. **Evidence that record-time reasoning is genuinely harder** (p=0.023 in mixed-effects regression controlling for question type, difficulty, and scenario), confirming the benchmark's diagnostic validity.

3. **Actionable diagnostic findings**: retrieval (not reasoning) is the primary failure bottleneck (∼55%); system confidence is anti-calibrated (AUROC < 0.5); memory backends are highly complementary (oracle +9.8pp over best single system).

4. **Cross-benchmark analysis** with LoCoMo (Raji et al., 2024) revealing that graph-based memory excels at adversarial robustness (97.5%) but struggles with temporal reasoning (37.5%).

2 Related Work

2.1 Temporal Reasoning Benchmarks

Recent years have seen a proliferation of temporal reasoning benchmarks for LLMs. TIME (Zhou et al., 2025) provides 38,522 questions across 11 subtasks at 3 difficulty levels, evaluating LLMs' ability to understand temporal expressions and reason about event sequences. ETRQA (Zhang et al., 2025) consolidates event temporal reasoning datasets for reading comprehension. TEST OF TIME (Kashyap et al., 2025) generates synthetic temporal reasoning tasks. MUSTQ (Chen et al., 2024) evaluates multilingual temporal KGQA across 10 languages. However, all treat time as a *single* dimension—none systematically distinguish event occurrence time from information recording time, a distinction critical for memory systems.

2.2 Memory System Benchmarks

The emergence of memory-augmented LLM agents has driven benchmark development. LoCoMo (Raji et al., 2024) (ACL 2024) evaluates long conver-

Benchmark	Year	Temporal	Dual TS	Mem Eval	Lang	Scale
LoCoMo	2024	Partial			EN	1,986
LongMemEval	2024				EN	4,500
TimE	2025				EN	38,522
ETRQA	2025				EN	Consolidated
Test of Time	2025				EN	Novel
MusTQ	2024				Multi	15K
BiTempQA	**2026**				ZH	608

Table 1: Comparison with related benchmarks. **Dual TS**: explicitly evaluates dual-timestamp reasoning. **Mem Eval**: tests memory system backends (not just LLMs).

sational memory with 1,986 questions across 5 types, including adversarial and temporal categories. LongMemEval (Wu et al., 2024) provides 4,500 questions testing 8 memory systems on single-hop, multi-hop, and temporal reasoning. MemoryAgentBench (Wang et al., 2025) focuses on multi-dataset agent memory evaluation. While these benchmarks include temporal questions, they do not systematically test *dual-timestamp* reasoning—the ability to distinguish "when this happened" from "when I learned about it."

2.3 Bitemporal Data Modeling

The distinction between event time (valid time) and record time (transaction time) is foundational in temporal database research (Snodgrass, 1999; Jensen and Snodgrass, 1999). Bitemporal data models track both dimensions, enabling queries like "what did the system know at time T_1 about events that occurred at time T_2?" This concept has been extensively studied in database systems but has not been applied to evaluate LLM memory systems. BiTempQA bridges this gap by bringing bitemporal reasoning into the memory system evaluation paradigm.

2.4 Positioning

Table 1 positions BiTempQA relative to existing work. BiTempQA is the *only* benchmark that (1) evaluates dual-timestamp reasoning and (2) tests multiple memory system backends.

3 BiTempQA Benchmark

3.1 Problem Formulation

Following bitemporal data modeling (Snodgrass, 1999), we define two temporal dimensions:

- **Event Time** (t_e): When an event actually occurred in the real world.
- **Record Time** (t_r): When the memory system recorded (learned about) the event.

A memory entry is a tuple $m = (\text{text}, t_e, t_r)$, and a memory store at logical time T contains all entries with $t_r \leq T$. A question may require reasoning about t_e only, t_r only, or both simultaneously.

3.2 Scenario Design

We design **10 scenario types** spanning 6 domains (corporate, academic, social, fictional, historical, scientific):

1. **Entity Attribute Evolution**: Attributes change over time (e.g., job titles).
2. **Relationship Evolution**: Social/professional relationships change.
3. **Contradictory Information**: Different sources report conflicting facts.
4. **Late-Arriving Facts**: Events recorded well after they occurred ($|t_r - t_e|$ large).
5. **Future-Dated Information**: Events recorded before their stated occurrence ($t_r < t_e$).
6. **Entity Identity Resolution**: Same entity referenced differently across time.
7. **Knowledge Retraction**: Previously recorded facts are later revoked.
8. **Multi-Source Information**: Multiple independent sources report overlapping events.
9. **Gradual Accumulation**: Knowledge builds incrementally over time.
10. **Temporal Ambiguity**: Unclear or ambiguous temporal references.

Each scenario consists of 3–10 memory writes, each annotated with explicit t_e and t_r ISO 8601 timestamps. Scenarios are designed so that $t_e \neq t_r$ in the majority of cases, creating natural bitemporal tension.

Temporal Requirement	Count (%)	Example Question Type
Both timestamps required	174 (56.5%)	"Did the system know X when it happened?"
Event-time only	110 (35.7%)	"When did event X occur?"
Record-time only	24 (7.8%)	"When was fact X first recorded?"
Version tracking needed	78 (25.3%)	"What is the current value of X?"
Knowledge retraction	28 (9.1%)	"Is the earlier claim about X still valid?"

Table 2: Dual-timestamp composition of BiTempQA. Over half of questions require reasoning about both t_e and t_r simultaneously. Categories are not mutually exclusive.

3.3 Question Taxonomy

Questions are organized into **9 types** at **3 difficulty levels**:

Level 1 (Easy): Point-in-time retrieval, temporal ordering, first-recorded detection.

Level 2 (Medium): Period queries, change detection, multi-hop temporal reasoning.

Level 3 (Hard): Counterfactual reasoning, version conflict resolution, complex temporal composition.

Questions use three answer formats: multiple choice (4 options, 43.8%), abstractive (free text, 53.0%), and boolean (3.2%). Every question is annotated with metadata flags: `requires_event_time`, `requires_record_time`, `requires_version_tracking`, and `requires_knowledge_retraction`.

3.4 Dataset Statistics

BiTempQA contains **308 QA pairs** across 112 scenarios. Table 2 shows the dual-timestamp composition.

All QA pairs are generated via Qwen2.5-72B-Instruct with structured prompts, then validated for parseability and temporal consistency. The dataset is Chinese-language, reflecting the growing importance of multilingual NLP evaluation.

3.5 Evaluation Pipeline

We implement a unified **Retrieve→Generate→Judge** pipeline (Figure 1):

Step 1—Retrieve: The memory backend retrieves the top-k context passages given the question and scenario memory store.

Step 2—Generate: DeepSeek-V3 (DeepSeek-AI et al., 2024) generates an answer conditioned on the retrieved context.

Step 3—Judge: An independent LLM call scores the generated answer against the gold answer. For MC questions, exact letter match is used (100% parse rate). For abstractive and boolean questions, DeepSeek-V3 in `answer_judge` mode evaluates semantic correctness with lenient matching. Without the LLM judge, abstractive accuracy drops to ∼10% (text matching too strict); with it, accuracy is ∼75% (fair semantic evaluation).

4 Experimental Setup

4.1 Memory Backends

We evaluate five memory backends spanning different retrieval paradigms:

- **FAISS** (Johnson et al., 2019): Dense vector similarity search using SiliconFlow's multilingual embedding API, top-k=5 retrieval.
- **BM25**: Sparse keyword matching via rank-BM25, top-k=5 retrieval.
- **Naive RAG**: Sequential passage concatenation (no relevance ranking), top-k=5.
- **Simple Knowledge Graph**: In-memory graph with entity/relation extraction, returning connected subgraphs (5–10 facts per query).
- **ChromaDB** (Chroma, 2023): Persistent vector database using its default embedding function (all-MiniLM-L6-v2).

All systems share identical chunking, top-k parameters, generation prompts, and the LLM backbone (DeepSeek-V3 via SiliconFlow API), ensuring differences reflect retrieval design choices.

4.2 Cross-Benchmark Evaluation

We additionally evaluate all backends (except ChromaDB) on the LoCoMo benchmark (Raji et al., 2024), which contains 1,986 MC questions across 10 long conversations (300–700 turns each). This enables cross-benchmark analysis of system behavior under different conditions: BiTempQA's short, dual-timestamp scenarios vs. LoCoMo's long, single-timestamp conversations.

4.3 Statistical Methods

We employ: (1) Bootstrap 95% confidence intervals (10,000 iterations); (2) McNemar's test with continuity correction for pairwise system comparisons; (3) Mixed-effects logistic regression with

System	Accuracy	95% CI	F1	Ver. Recall
FAISS	**75.6%**	[70.8, 80.2]	**0.870**	**79.5%**
Naive RAG	74.0%	[69.2, 78.9]	0.867	74.4%
BM25	73.4%	[68.2, 78.2]	0.867	69.2%
Simple KG	67.2%	[61.7, 72.4]	0.841	75.6%
ChromaDB†	41.2%	[35.7, 46.8]	0.744	37.2%

Table 3: Overall results on BiTempQA (308 questions). † ChromaDB uses a different embedding model (see §7).

System	Event-Only (110Q)	Record-Only (24Q)	Both Req. (174Q)	Version Track. (78Q)
FAISS	79.1%	66.7%	**74.7%**	**79.5%**
Naive RAG	78.2%	62.5%	73.0%	74.4%
BM25	77.3%	66.7%	71.8%	69.2%
Simple KG	66.4%	58.3%	69.0%	75.6%
ChromaDB	44.5%	50.0%	37.9%	37.2%
Average	**69.1%**	60.8%	65.3%	

Table 4: Accuracy by temporal requirement. Record-only questions (n=24) are hardest across all systems; the gap is suggestive given the small sample size.

scenario random intercepts to control for question type, difficulty, and temporal requirements; (4) LLM Judge cross-validation (DeepSeek-V3 vs. Qwen2.5-72B, Cohen's κ=0.455).

5 Results

5.1 Overall Performance

Table 3 presents overall accuracy by system. FAISS leads at 75.6%, followed closely by Naive RAG (74.0%) and BM25 (73.4%). Bootstrap 95% CIs overlap substantially among the top three, and pairwise McNemar's tests show no significant differences (p>0.3). Simple KG (67.2%) and ChromaDB (41.2%) are significantly worse (p<0.01 and p<0.001 respectively). ChromaDB uses a different default embedding model (all-MiniLM-L6-v2) than the other vector-based systems; we include it for completeness but note this confound.

5.2 Record-Time Reasoning Is Genuinely Harder

Table 4 shows accuracy by temporal requirement. Record-only questions average 60.8% vs. 69.1% for event-only—an 8.3pp gap consistent across all five systems. To confirm this is not a distribution artifact, we fit a mixed-effects logistic regression with scenario random intercepts, controlling for system identity, question type, difficulty level, and all temporal requirement flags.

Predictor	Coef.	z	p
System effects (ref: FAISS)			
ChromaDB	−0.344	−10.53	<0.001
Simple KG	−0.084	−2.58	0.010
BM25	−0.023	−0.70	0.487
Naive RAG	−0.016	−0.50	0.619
Temporal requirements			
Record-time required	−0.079	−2.26	**0.023**
Event-time required	0.061	1.08	0.280
Version tracking	0.050	1.29	0.199
Knowledge retraction	−0.080	−1.50	0.135
Scenario variance		σ^2=0.049	

Table 5: Mixed-effects logistic regression (scenario random intercepts, N=1,540). Record-time requirement is the only temporal factor reaching significance.

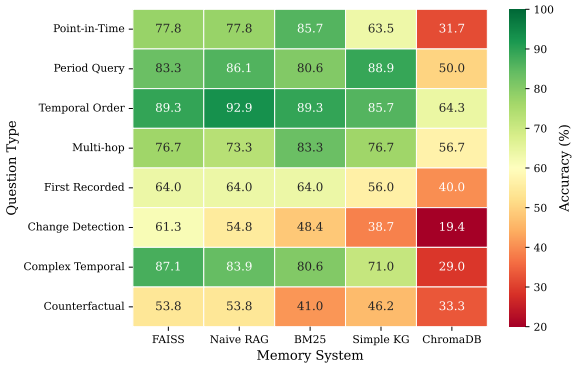

Figure 2: System × question type accuracy heatmap. No single system dominates all types; the 40.6pp spread across types confirms the typology captures distinct challenges.

The record-time coefficient (β=−0.079, p=0.023) remains significant after controlling for all confounds (Table 5), confirming that the dual-timestamp dimension captures a genuinely harder reasoning capability. Notably, difficulty levels (L2/L3) are *not* independently predictive after controls (p=0.103, p=0.278), suggesting the difficulty gradient is largely mediated by the distribution of harder question types at higher levels.

5.3 Per-Question-Type Analysis

Figure 2 reveals substantial system×type interactions. FAISS leads on complex temporal (87.1%) and counterfactual (53.8%); BM25 excels at point-in-time (85.7%) and multi-hop (83.3%); Simple KG tops period query (88.9%). The 40.6pp spread across question types (counterfactual: 48.7% to temporal order: 89.3%) validates the typology.

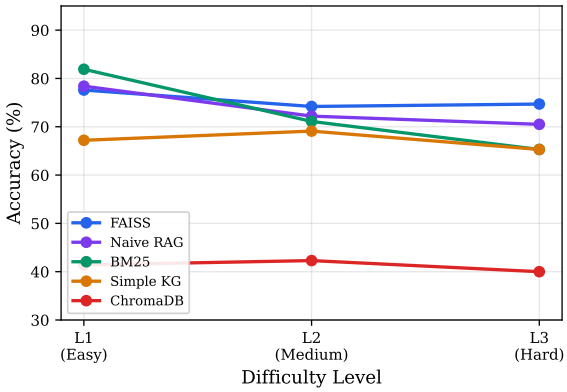

Figure 3: Accuracy degradation from L1 to L3. FAISS is most robust (−2.9pp); BM25 is most fragile (−16.6pp).

5.4 Difficulty Calibration

All five systems show monotonic L1→L3 degradation, but with strikingly different robustness: FAISS drops only 2.9pp vs. BM25's 16.6pp (Figure 3). FAISS's dense representations appear to maintain semantic relevance even under complex compositional demands, while BM25's keyword matching deteriorates sharply when query–document lexical overlap decreases.

5.5 System Complementarity

An oracle ensemble (correct if *any* system answers correctly) achieves **85.4%**—a 9.8pp gain over the best single system—confirming that memory backends have substantially non-overlapping error profiles (Figure 4). FAISS-ChromaDB Jaccard overlap is only 0.446 (lowest), while FAISS-SimpleKG is 0.746. Simple KG uniquely solves 6 items that FAISS misses, and ChromaDB uniquely solves 6—despite its low overall accuracy. 45 items (14.6%) are answered incorrectly by *all* systems, establishing a hard irreducible ceiling.

5.6 Cross-Benchmark Validation with LoCoMo

Table 6 compares BiTempQA results with LoCoMo evaluation. The most striking finding is Simple KG's asymmetry: 97.5% adversarial accuracy on LoCoMo (best of all systems, +34pp above average) but 37.5% temporal reasoning on LoCoMo (worst, −9pp) and 67.2% on BiTempQA. This reflects a fundamental trade-off: graph structures resist adversarial hallucination (returning empty context for non-existent facts) but lack the retrieval breadth needed for temporal reasoning.

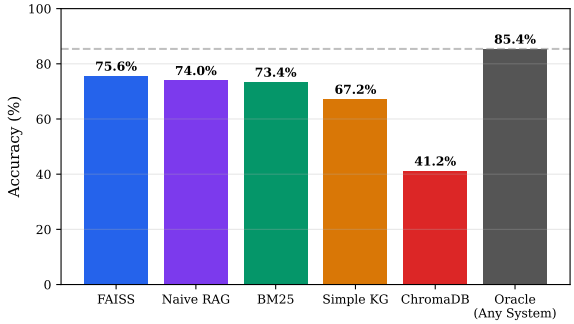

Figure 4: Oracle ensemble reaches 85.4% vs. 75.6% best single system. 45 irreducible failures suggest retrieval-independent challenges.

	BiTempQA		LoCoMo		
System	Overall	Temporal	Overall	Temporal	Adversarial
FAISS	**75.6%**	**75.6%**	**57.0%**	**46.9%**	63.0%
Naive RAG	74.0%	74.0%	54.4%	42.7%	69.7%
BM25	73.4%	73.4%	55.8%	43.8%	64.8%
Simple KG	67.2%	67.2%	56.8%	37.5%	**97.5%**

Table 6: Cross-benchmark comparison. Simple KG shows extreme adversarial-temporal asymmetry.

6 Diagnostic Analysis

6.1 Retrieval Is the Bottleneck

We sampled 20 FAISS failures (best system) and classified them by retrieval context length. All 20 had short or empty retrieval context (<200 characters), compared to ∼500 chars average for correct answers. Heuristic classification estimates ∼55% retrieval failure vs. ∼45% reasoning failure, suggesting that improving retrieval quality—particularly temporal-aware retrieval—would yield larger gains than improving reasoning capability.

6.2 Confidence Is Anti-Calibrated

A striking diagnostic finding: vector-based systems produce *anti-calibrated* confidence scores. Table 7 shows Expected Calibration Error (ECE), Brier score, and AUROC for each system.

FAISS's AUROC of 0.453 means confidence scores are *worse than random* at predicting correctness: wrong answers receive higher confidence than correct ones. This has practical implications: agents cannot rely on internal confidence to decide whether to trust a temporal memory answer or seek verification.

System	ECE↓	Brier↓	AUROC
FAISS	0.163	0.229	0.453
Naive RAG	0.319	0.303	0.552
BM25	0.235	0.769	0.467
Simple KG	0.672	0.672	0.500
ChromaDB	0.412	0.412	0.500

Table 7: Confidence calibration metrics. FAISS and BM25 show AUROC < 0.5—wrong answers receive *higher* confidence than correct ones.

System	Retraction (28Q)	Version Track. (58Q)	Neither (232Q)
FAISS	67.9%	**82.8%**	74.8%
Naive RAG	67.9%	75.9%	74.3%
BM25	57.1%	74.1%	**75.2%**
Simple KG	53.6%	81.0%	65.5%
ChromaDB	39.3%	37.9%	42.3%

Table 8: Accuracy by knowledge dynamics type. Simple KG: strong version tracking but weak retraction. Bootstrap CIs for retraction span 35pp due to small n=28.

6.3 Knowledge Retraction vs. Version Tracking

We decompose performance on questions involving knowledge dynamics (Table 8). Simple KG shows an asymmetric profile: strong at version tracking (81.0%, close to FAISS's 82.8%) but weak at knowledge retraction (53.6%). This suggests graph structures naturally support tracking fact evolution but struggle with explicitly revoking superseded knowledge—a distinction with implications for memory architecture design.

6.4 Implications for Memory System Design

Our findings suggest three design recommendations:

1. Temporal-aware retrieval. Since $\sim$55% of failures are retrieval-limited, adding temporal metadata filtering (e.g., "retrieve only facts recorded before T") or bitemporal indexing would likely yield larger gains than improving the reasoning model.

2. Hybrid vector+graph architectures. The complementarity between systems (oracle +9.8pp) and Simple KG's unique adversarial robustness suggest combining dense retrieval's breadth with graph structure's precision would outperform either alone.

3. External verification for temporal claims. Anti-calibrated confidence means internal signals cannot identify uncertain temporal answers. External verification (e.g., cross-checking against source timestamps) should be triggered for all bitemporal queries.

7 Limitations

We acknowledge several limitations:

Scale and language. BiTempQA contains 308 items in Chinese only. While sufficient for diagnostic analysis (bootstrap CIs $\leq$8pp for major categories), small subcategories (record-only n=24, retraction (n=28) have wide CIs that limit strong quantitative claims. Expanding to 600–1000 items and adding English translations are priorities for future work.

LLM-generated data. All QA pairs are generated by Qwen2.5-72B-Instruct and have not undergone human expert validation. We mitigate this through: (1) structured generation with strict output schemas; (2) automated consistency checks; (3) LLM Judge cross-validation (κ=0.455). However, some items may contain ambiguous or incorrectly specified answers. Human validation on a stratified sample (100–150 items, 2 native annotators) is the single highest-priority next step.

ChromaDB confound. ChromaDB uses a different default embedding model (all-MiniLM-L6-v2, 22M parameters) than the other vector systems (SiliconFlow API, multilingual). Its poor performance (41.2%) likely reflects this configuration difference rather than fundamental limitations of persistent vector databases. We include ChromaDB for completeness but exclude it from primary comparative claims.

Backend scope. We evaluate memory storage/retrieval backends rather than full agent memory systems (e.g., Mem0, Zep, Graphiti). These production systems incorporate additional components (memory consolidation, summarization, conflict resolution) that may improve bitemporal reasoning. Evaluating such systems is ongoing work.

LLM Judge leniency. DeepSeek-V3 is more lenient than Qwen2.5-72B (80% vs. 60% accuracy on a 50-item sample), with 11/12 disagreements favoring acceptance. Reported accuracies may be slightly inflated relative to a stricter judge, though relative rankings are preserved.

8 Conclusion

We presented BiTempQA, the first diagnostic benchmark explicitly designed to evaluate bitem-

poral reasoning—reasoning about *when events occurred* vs. *when the system learned about them*—in LLM agent memory systems. Our evaluation of five memory backends reveals three key findings:

1. **Record-time reasoning is genuinely harder** than event-time reasoning (p=0.023), confirming that the dual-timestamp distinction captures a real cognitive challenge.

2. **Retrieval, not reasoning, is the primary bottleneck** ($\sim$55% of failures), suggesting that temporal-aware retrieval mechanisms would yield the largest improvements.

3. **Memory backends are highly complementary**: an oracle ensemble achieves 85.4% (+9.8pp), and graph-based memory offers unique adversarial robustness (97.5%) despite lower temporal accuracy.

These findings provide actionable guidance for building temporally-aware memory systems: invest in bitemporal indexing, combine vector breadth with graph precision, and use external verification rather than internal confidence.

Future work includes: human expert validation, English translation, evaluation of production memory systems (Mem0, Graphiti, TMG), and expansion to longer scenarios with richer temporal dynamics.

References

Xinyin Chen and 1 others. 2024. MusTQ: Multilingual temporal question answering over knowledge graphs. In *ACL Findings*.

Chroma. 2023. Chromadb: The ai-native open-source embedding database. *https://www.trychroma.com/*.

DeepSeek-AI and 1 others. 2024. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*.

Christian S. Jensen and Richard T. Snodgrass. 1999. Temporal data management. *IEEE Transactions on Knowledge and Data Engineering*, 11(1):36–44.

Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2019. FAISS: A library for efficient similarity search and clustering of dense vectors. *arXiv preprint arXiv:2401.08281*.

Aditya Kashyap and 1 others. 2025. Test of time: Benchmarking LLM temporal reasoning. In *International Conference on Learning Representations (ICLR)*.

Inioluwa Deborah Raji and 1 others. 2024. Locomo: Evaluating LONG-Context LLM memorization. In *Proceedings of ACL 2024*.

Richard T. Snodgrass. 1999. TSQL2: A temporal query language. *ACM Transactions on Database Systems*, 24(1):1–58.

Yu Wang and 1 others. 2025. MemoryAgentBench: Evaluating LLM agents on memory tasks. *arXiv preprint*.

Di Wu and 1 others. 2024. LongMemEval: Benchmarking chat-optimized LLM memory systems. In *Proceedings of ACL 2024*.

Song Zhang and 1 others. 2025. ETRQA: Event temporal reasoning question answering. In *ACL Findings*.

Zhi Zhou and 1 others. 2025. TimE: Temporal reasoning of events in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.

A Appendix

A.1 Full Per-Question-Type Results

Table 9 provides complete per-question-type accuracy across all systems and types.

A.2 LLM Judge Cross-Validation Details

We cross-validated DeepSeek-V3 against Qwen2.5-72B-Instruct on 50 abstractive QA pairs from the FAISS system (random sample, seed=42).

Question Type	Level	FAISS	BM25	Naive RAG	Simple KG	ChromaDB
Point-in-Time	L1	77.8%	**85.7%**	77.8%	63.5%	31.7%
Temporal Order	L1	89.3%	89.3%	**92.9%**	85.7%	64.3%
First Recorded	L1	64.0%	64.0%	64.0%	56.0%	40.0%
Period Query	L2	83.3%	80.6%	**86.1%**	88.9%	50.0%
Change Detection	L2	**61.3%**	48.4%	54.8%	38.7%	19.4%
Multi-hop Temporal	L2	76.7%	**83.3%**	73.3%	76.7%	56.7%
Counterfactual	L3	**53.8%**	41.0%	53.8%	46.2%	33.3%
Complex Temporal	L3	**87.1%**	80.6%	83.9%	71.0%	29.0%
Version Conflict	L3	—	—	—	—	—
Average		74.2%	71.5%	72.2%	65.8%	40.7%
Range		48.7% (Counterfactual) to 92.9% (Temporal Order): 44.2pp spread				

Table 9: Complete per-question-type accuracy. Version Conflict excluded due to insufficient samples.

Metric	Value
Agreement rate	76.0% (38/50)
Cohen's κ	0.455 (moderate)
DeepSeek-V3 judged correct	80.0% (40/50)
Qwen2.5-72B judged correct	60.0% (30/50)
Disagreement breakdown (12 items)	
DeepSeek lenient	11/12 (91.7%)
Qwen lenient	1/12 (8.3%)
Pattern: partial answer	6/12
Pattern: answer too short	6/12

Table 10: LLM Judge cross-validation. DeepSeek-V3 is systematically more lenient, accepting partial answers that Qwen rejects.